

[CoRL'24] Robot See Robot Do: Imitating Articulated Object Manipulation with Monocular 4D Reconstruction

1. Link: <https://robot-see-robot-do.github.io/>
2. Authors and institution: Justin Kerr*, Chung Min Kim*, Mingxuan Wu, Brent Yi, Qianqian Wang, Ken Goldberg, Angjoo Kanazawa from UCB
 1. Angjoo Kanazawa is an AP in BAIR, UCB, she also holds position in Google research team, advisory board of Wonder Dynamics <https://wonderdynamics.com/> and LumaAI <https://lumalabs.ai/dream-machine>.

TL;DR Robot See Robot Do uses a 4D Differentiable Part Model (4D-DPM) to visually imitate articulated motions from an object scan and single monocular video.

TODOs

1. TRY Nerfstudio <https://arxiv.org/pdf/2302.04264> WITH QINGYUAN DATA!!!

Miscellaneous work

1. [CVPR'24] GARField, a previous work by the team <https://www.garfield.studio/>. GARField hierarchically groups in 3D by training a scale-conditioned affinity field from multi-level masks.

Thoughts and criticisms

1. A key insight that enables visual imitation across a morphology gap is not to focus on the exact motion of the manipulator, but **observe the consequence of the action at the object level**, so one's goal is planning to replicate the perceived 3D motion of object.
2. The detection of part of interest/part of task is computed via trajectory of human hands, which sometimes is not attainable

- for new tasks. Can we leverage VLM to solve this?
3. The way authors use 3DGS does not consider all physical parameters, though they demonstrate manipulation of deformable objects. That comes to the question "at what level of difficulty for a non-prehensile manipulation task, an estimation of physical parameters is necessary?"
 4. The experiment stops at prehensile manipulation level.

Related works

1. **Recovering 3D Motion for Objects with Moving Parts**

1. definition: Reconstructing the 3D motion trajectory of articulated object from a single video
2. Inputs:
 1. sequences of articulated motions of objects
 2. single point clouds
3. requirements
 1. RGBD videos/multi-view observations
4. known kinematic chains/category-specific priors

2. **Learning from one demonstration.**

1. use human data
 1. need additional robot teleoperation data for target tasks
 2. paired human-robot data to bridge the morphology gap
2. learn policy from single demonstration
 1. require extensive in-domain data
 2. require well-curated meta-training tasks for training

3. **Object-centric representations for robot manipulation**

1. methods:
 1. representation of specific classes of data
 1. using 3D structure as a signal for contrastive learning
 2. intermediate object properties as a bottleneck for image prediction
 3. canonical object views as conditioning
2. generalization methods
 1. imitation
 2. RL
 3. semantic

- 4. relational
- 5. uncertainty-based
- 6. symbolic reasoning
- 7. pair-wise reasoning

4. Feature fields for robotics

- 1. NERF
- 2. 3DGS
- 3. Feature3DGS

Contributions

- 1. A method for recovering 3D part motion from a monocular video with differentiable rendering.
- 2. A manipulation framework that leverage the object-centered representation

Key concepts

Task definition

Assumptions internal part deformation is negligible **Robot**: bimanual robot with parallel-jaw grippers **Object**: an object with two or more movable parts **Observation**:

- 1. a multi-view object scan of an object
- 2. a monocular human demonstration

Goal: manipulate the object through the same configuration change starting from an unknown location in the robot's workspace.

Algorithm

Constructing 4D Differentiable Part Models、

- 1. construct a 3DGS model of object
- 2. construct a GARField model of object, cluster the gaussians into subgroups
- 3. train a Feature 3DGS model o each subgroup

Monocular 3D Part Motion Recovery

Analysis-by-synthesis approach we optimize, or synthesize, a model of the object parts over time, to understand (i.e analyze)

their motion

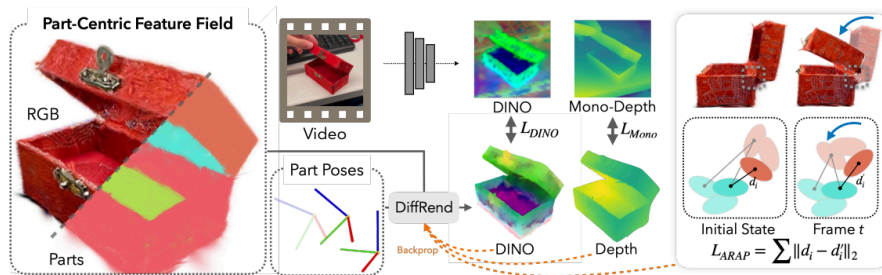


Figure 3: **4D Differentiable Part Models (4D-DPM)**. *Left*: DINO features and depth are rendered from per-timestep optimizable part pose parameters, and compared with extracted DINO features and monocular depth from the input frame. *Right*: an ARAP loss penalizes gaussians from deviating too far from their initial configuration with respect to neighbors. Together these losses flow backwards into the part poses and are optimized with gradient descent to recover 3D part motion.

1. 4D part pose ---> trajectory of SE3 pose
 2. backpropagate pixel error into 3D pose delta
 3. better than feed-forward as it can be intergrated with geometry rigidity and temporal smoothing through optimization **Part Motion Optimization from Video**
 4. get rendered RGBD image and feature maps
 5. optimize pose of the part by B.P. tricks:
 1. blur them with a kernel at ViT size
 2. clip features with low alpha in 3DGS
 3. 1st done per frame, them optimize all in once with Laplacian Temporal smoothness **3D Regularization Priors** two additional losses are added into DINO loss
 4. depth loss from DepthAny
 5. rigidity loss
 1. find boundaries of each gaussian parts
 2. derive a distance between neighboring pairs
 3. add it to the loss
- $$\sum_{ij} \rho(d_{init}^{ij} - d_{current}^{ij})$$
- Initialization**
6. initialize the object pose
 1. get location in 2D image
 1. compare object 3D DINO feature with 2D image feature
 2. find mutual nearest neighbors btw 3DGS and pixel feature
 3. pixel centroids --> ray
 7. after that, the procedure is used in stereo depth instead of monocular depth

Object Motion and Grasp Planning

Hand-Guided Part Selection Decide which part to move, this part firstly estimate the pose of human hands in the view by previous work, then estimate each part-hand distance and rank them, finally select the part with the highest rank.

Part-Centric Grasp Planning sample 20 antipodal grasp axes per part use analytical methods then augment them to 480 grasps

Robot Trajectory Planning

1. use trajectory optimization methods to a list of robot end-effector pose motions
2. use cuRobo for collision avoidance checks and to plan robot approach motions, execute the first successful trajectory

Experiments

Hardware

1. ABB YuMi bimanual arms
2. soft 3D-printed parallel-jaw grippers
3. ZED 2 stereo camera

4. Polycam phone scanner




Figure 6: Example Robot Executions. Arrows indicate direction of motion.